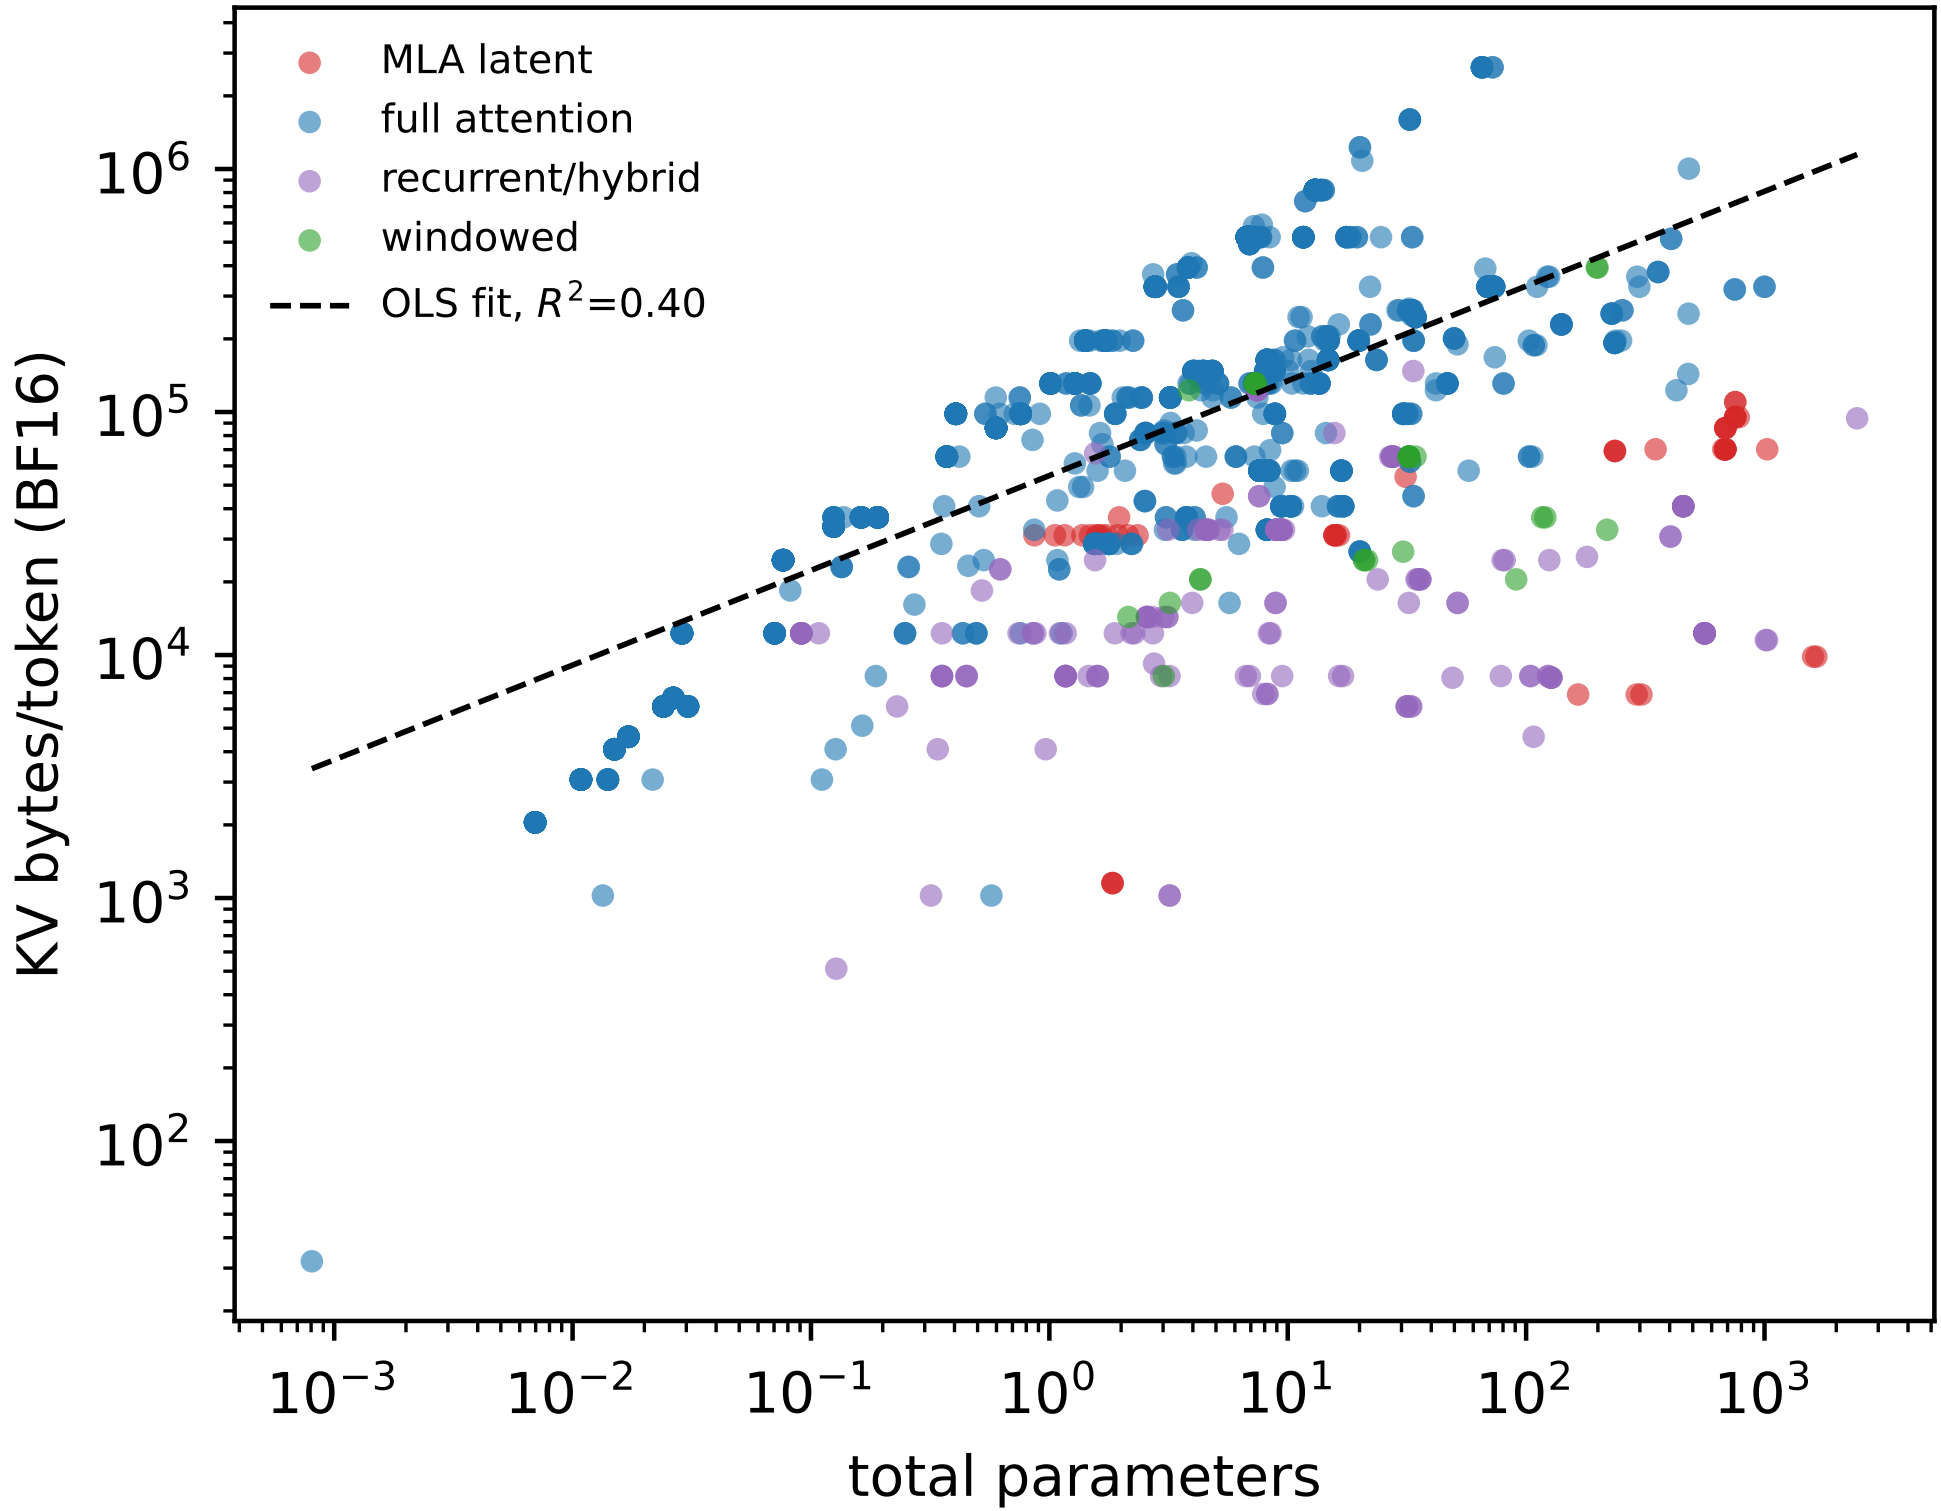
MLA latent
full attention
recurrent/hybrid
windowed
OLS fit, R^2=0.40
KV bytes/token (BF16)
10^6
10^5
10^4
10^3
10^2
10^{-3}
10^{-2}
10^{-1}
10^0
10^1
10^2
10^3
total parameters