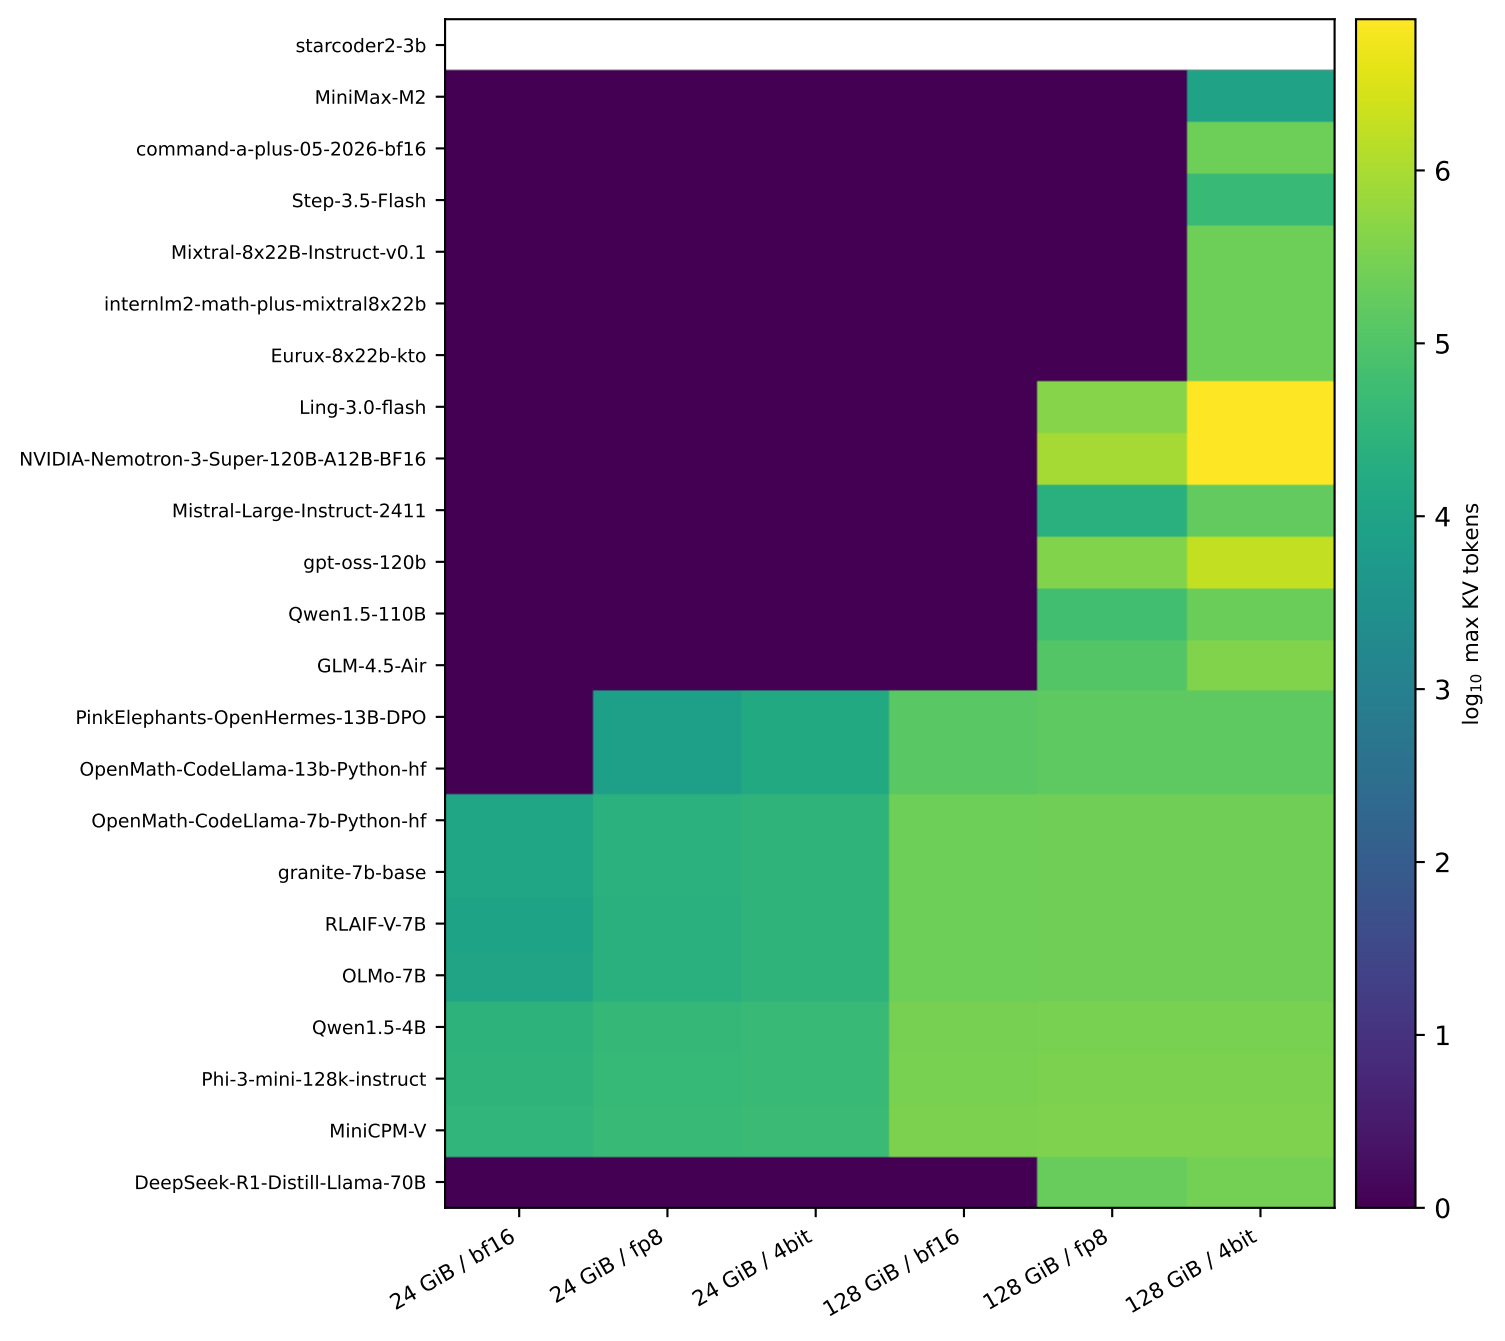
starcoder2-3b
MiniMax-M2
command-a-plus-05-2026-bf16
Step-3.5-Flash
Mixtral-8x22B-Instruct-v0.1
internlm2-math-plus-mixtral8x22b
Eurux-8x22b-kto
Ling-3.0-flash
NVIDIA-Nemotron-3-Super-120B-A12B-BF16
Mistral-Large-Instruct-2411
gpt-oss-120b
Qwen1.5-110B
GLM-4.5-Air
PinkElephants-OpenHermes-13B-DPO
OpenMath-CodeLlama-13b-Python-hf
OpenMath-CodeLlama-7b-Python-hf
granite-7b-base
RLAIF-V-7B
OLMo-7B
Qwen1.5-4B
Phi-3-mini-128k-instruct
MiniCPM-V
DeepSeek-R1-Distill-Llama-70B
24 GiB / bf16
24 GiB / fp8
24 GiB / 4bit
128 GiB / bf16
128 GiB / fp8
128 GiB / 4bit
log10 max KV tokens
0
1
2
3
4
5
6